



ECS708P - MACHINE LEARNING - 2017/18



ASSIGNMENT 1 (LINEAR REGRESSION + LOGISTIC REGRESSION)

Table of Contents

Table of Contents.....	1
Assignment 1: Part 1 - Linear Regression.....	3
1- Linear Regression with One Variable	3
Task 1:	3
2. Linear Regression with Multiple Variables.....	6
Task 2:	6
How much does your algorithm predicts that a house with 1650 sq. ft. and 3 bedrooms cost?.....	8
How about 3000 sq. ft. and 4 bedrooms?	8
3. Regularized Linear Regression	9
Task 3:	9
Assignment 1 – Part 2: Logistic Regression and Neural Networks	12
Task 1:	12
Task 2:	13
1.1. Cost function and gradient for logistic regression.....	13
Task 3:	13
Task 4:	15
Task 5:	16
Task 6:	16
Task 7:	18
Task 8:	19
Task 9:	20
2. Neural Network.....	21
Task 10.	21
Task 11.	23

In this exercise, we have one variable, X , and we building a model to predict one

Hypothesis and Cost Graph are flat as the hypothesis is not calculated yet.



Assignment 1: Part 1 - Linear Regression

1- Linear Regression with One Variable

Task 1:

Modify the function *calculate_hypothesis.m* to return the predicted value for a single specified training example. Include in the report the corresponding lines from your code.


Calculate Hypothesis:

```
function hypothesis = calculate_hypothesis (X, theta,  
training_example)  
  
%CALCULATE_HYPOTHESIS This calculates the hypothesis for a given X,  
%theta and specified training example  
  
hypothesis = X(training_example,1) *theta (1) +  
X(training_example,2) *theta (2);  
  
end
```


After making changes to the hypothesis and adding this line of code

```
hypothesis = X(training_example,1) *theta (1) +  
X(training_example,2) *theta (2);
```

The hypothesis graph below shows the changes and not the flat line anymore.





When now we run mllab1.m we can see the cost going down overtime.



Now we modified the values for the learning rate, alpha in mllab1.m and the results are posted below.



First, we will increase the value of alpha to 0.90.

`alpha = 0.90;`



Secondly, we will decrease the value to 0.001.

`alpha = 0.001;`



2. Linear Regression with Multiple Variables

Task 2:

Modify the functions *calculate_hypothesis* and *gradient_descent* to support the new hypothesis function. This should be sufficiently general so that we can have any number of extra variables. Include the relevant lines of the code in your report.

Following is the *calculate_hypothesis*, a few lines of code were added as shown below.

```
function hypothesis = calculate_hypothesis (X, theta,
training_example)

%CALCULATE_HYPOTHESIS This calculates the hypothesis for a given X,
%theta and specified training example

hypothesis = sum (X (training_example, :) * theta');

end
```

After making changes to the hypothesis and adding this line of code

```
hypothesis = sum (X (training_example, :) * theta');
```

Following is the code that was added to *gradient_descent* to support the new hypothesis function, this was added to get the output for $\theta_0 = \theta(1)$; $\theta_1 = \theta(2)$; $\theta_2 = \theta(3)$;

```
%update theta(1) and store in temporary variable theta_0


sigma = 0.0;
for i = 1:m
hypothesis = X(i, 1) * theta(1) + X(i, 2) * theta(2);
output = y(i);
sigma = sigma + (hypothesis - output);
end

theta_0 = theta_0 - ((alpha * 1.0) / m) * sigma;
```

The theta was then updates as below:


```
theta = [theta_0, theta_1, theta_2];
```

Following is the graph result for the hypothesis (The price has fallen over time)




Now we will see how different values of alpha affect the convergence of the algorithm. The results are as posted below.

`alpha = 0.90;`




```
alpha = 0.001;
```



How much does your algorithm predicts that a house with 1650 sq. ft. and 3 bedrooms cost?

We use the formula for normalization and came up with the following line of code.

```
p_1 = (1650 - mean_vec(1))/std_vec(1);  
r_1 = (3 - mean_vec(1))/std_vec(1);
```

To predict the value, we are using the following formula which will generate the results.

```
q_1 = t(1)+p_1*t(2)+r_1*t(2);
```

The results are as follows:

2.9758e+04

How about 3000 sq. ft. and 4 bedrooms?

We use the formula for normalization and came up with the following line of code.

```
a_1 = (3000 - mean_vec(1))/std_vec(1);  
b_1 = (4 - mean_vec(1))/std_vec(1);
```

To predict the value, we are using the following formula which will generate the results.

```
z_1 = t(1)+a_1*t(2)+b_1*t(2);
```

The results are as follows:

2.0847e+05

3. Regularized Linear Regression

Task 3:

Note that the punishment for having more terms is not applied to the bias. This cost function has been implemented already in the function `compute_cost_regularised`. Modify `gradient_descent` to use the `compute_cost_regularised` method instead of `compute_cost`. Include the relevant lines of the code in your report and a brief explanation.

Modifying *gradient_descent* to create:

```
theta_0 = theta(1);  
theta_1 = theta(2);  
theta_2 = theta(3);  
theta_3 = theta(4);  
theta_4 = theta(5);  
theta_5 = theta(6);
```

Update the theta for the new theta values.

%update theta

```
theta = [theta_0, theta_1, theta_2, theta_3, theta_4, theta_5];
```

Update the cost_vector to use `compute_cost_regularised` instead of `compute_cost`

%update cost_vector

```
cost_vector = [cost_vector; compute_cost_regularised(X, y, theta, l)];
```

Calculate_hypothesis is modified to run mllab3.m



```
function hypothesis = calculate_hypothesis(X, theta, training_example)

%CALCULATE_HYPOTHESIS This calculates the hypothesis for a given X,
%theta and specified training example

hypothesis = sum(X(training_example,:)* theta');



end
```

Plot the hypothesis function found at the end of the optimization.



After a few attempts to find the best value for alpha, I am came up with 0.007

The results are as posted below.




Assignment 1 – Part 2: Logistic Regression and Neural Networks

Task 1:

Include in your report the relevant lines of code and the result of the running the `plot_sigmoid_function.m`.


plot_sigmoid_function.m

```
h=plot_sigmoid;
```



Plot the data again to see what it looks like in this new format. Enclose this in your report.

A scatter plot showing the relationship between x_1 (horizontal axis) and x_2 (vertical axis) for two classes. The horizontal axis ranges from -3 to 3, and the vertical axis ranges from -3 to 2. The blue circles are clustered in the upper right region, while the red crosses are clustered in the lower left region. There is a significant overlap between the two classes in the central region of the plot.



Task 3:

Modifying the calculate_hypothesis: to:

```
hypothesis = 0.0;
```

```
%Calculate the hypothesis for the i-th training example in X.
```

%%%%%%%%%

```
result=sigmoid(hypothesis);
```

%END OF FUNCTION

Modifying the sigmoid.m

```
function output=sigmoid(z)
```

```
% modify this to return z passed through the sigmoid function
```


```
%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
```

```
output = zeros(size(z));
```

```
output = 1 ./ (1 + exp(-z));
```

```
%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
```

%END OF FUNCTION



Task 4:

Modify the line: `cost = 0.0` in `compute_cost(X,y,theta` so that it uses the cost function:


```
for i=1:m
hypothesis = calculate_hypothesis(X,theta,i);
output = y(i);

cost = -output*log(hypothesis)-(1-output)*log(1-hypothesis);

J = 1 / m * sum(-y .* log(hypothesis) - (1 - y) .* log(1 - hypothesis));

% modify this to calculate the cost function, using hypothesis and output
cost = 0.0;
J =J+cost;



end
```



Task 5:

Plot the decision boundary. This corresponds to the line where $\theta^T \mathbf{x} = 0$. For example, if $h\theta = \theta_1 x_1 + \theta_2 x_2$ the boundary is where $\theta_1 x_1 + \theta_2 x_2 = 0$



```
% plot our data and decision boundary
plot_data_function(X,y)
plot_boundary(X,t)
```





Task 6:

Run the code in lab2_lr_ex2.m several times. What is the general difference between the training and test error? When does the training set generalize well? Demonstrate two splits with good and bad generalisation and put both graphs in your report.



Training error: 0.069635



Training error:0.18386



Training error:0.072579



Task 7:

Run logistic regression on this dataset. How does the error compare to using the original features (i.e. the error found in Task 4)? Include in your report the error and an explanation on what happens


```
% after normalising we add the bias
X=[ones(size(X,1),1),X];
theta=ones(1,6);

% for question 7, modify the dataset X to have more features (in each row)
% append to X(i), the following features:

% here append x_2 * x_3 (remember that x_1 is the bias)
X(:,4) = X(:,2).*X(:,3);

% here append x_2 * x_2 (remember that x_1 is the bias)
X(:,5) = X(:,2).^2;

%here append x_3 * x_3 (remember that x_1 is the bias)
%X(:,6) = X(:,3).*X(:,3);
X(:,6) = X(:,3).^3;
```




Task 8:

In lab2_lr_ex3b.m the data is split into a test set and a training set. Add your new features from the question above. Modify the function `gradient_descent_training()` to store the current cost for the training set and testing set. Store the cost of the training set to `cost_array_training` and for the test set to `cost_array_test`


Gradient_descent_training

```
% update cost_array
% add code here: to update cost_array_training and cost_array_test
cost_array_training(it)=compute_cost(X, y, theta);
cost_array_test(it)= compute_cost(test_X,test_y,theta);
```


Data Size: 40, Training:0.019723, Test:0.018076



Data Size: 60, Training: 0.015277, Test: 0.046065



Data Size: 10, Training: 0.11859, Test: 0.037704



Task 9:

With the aid of a diagram of the decision space, explain why a logistic regression unit cannot solve the XOR classification problem.

To solve the problem, we need to introduce a new layer into our neural networks. This layer, often called the 'hidden layer', allows the network to create and maintain internal representations of the input. Here's is a network with a hidden layer that will produce the XOR truth table above: XOR Network.ⁱ

2. Neural Network

Task 10.

Implement backpropagation. Although XOR only has one output, this should support outputs of any size. Do this following the steps below.

Modified Sigmoid.m

```
function sigmoid_output=sigmoid(z)
% change this to apply the sigmoid to the data below:

sigmoid_output = 1.0 ./ (1.0 + exp(-z));

%sigmoid_output = 0.0;
end
```

Modified NeuralNetwork.m

```
% Step 1. Output deltas (used to change weights from hidden --> output)
output_deltas = zeros(1,length(nn.output_neurons));
outputs=nn.output_neurons;
for i=1:length(outputs)
    output_deltas(i) = (outputs(i) -
targets(i))*sigmoid_derivative(outputs(i));
end
% Step 2. Hidden deltas (used to change weights from input -->
output).
hidden_deltas = zeros(1,length(nn.hidden_neurons));
% hint... create a for loop here to iterate over the hidden
neurons and for each

##### my coe is in between
#####
% hidden neuron create another for loop to iterate over the
ouput neurons

%for i=1:length(nn.hidden_neurons)
%hiddens = nn.hidden_neurons;
%for j=1:length(hiddens)
%hidden_deltas =
sigmoid_derivative*(inputs(j)+(inputs(j))*(outputs));
%end
%end
##### my coe is in between
#####

% Step 3. update weights output --> hidden
for i=1:length(nn.hidden_neurons)
    for j=1:length(output_deltas)
        nn.output_weights(i,j) =nn.output_weights(i,j) -
(output_deltas(j) * nn.hidden_neurons(i) * learning_rate);
    end
end
```

```

        % here we are removing the bias from the hidden neurons as
there is no
        % connection to it from the layer below
hidden_deltas = hidden_deltas(2:end);


        % Step 4. update weights input --> hidden.
        % hint, use a similar process to step 3, except iterate over
the input neurons and hidden deltas

        %for i=1:length(nn.hidden_neurons)
            for i=1:length(output_deltas)
                nn.output_weights(i) =nn.output_weights(i) -
(output_deltas(i) * nn.hidden_neurons(i) * learning_rate);
            end
        %end

        % this is our cost function
J = 0.0;
        for t =1:length(targets)
            J = J + 0.5*(nn.output_neurons(t)-targets(t))^2;
        end
    end

    function propagation=forward_propagate(inputs,nn)
        %
        %Calculates the output by feeding the inputs forward through
the network
        %
        nn.reset_activations()
        % activate hiddens_neurons
        for i=1:length(nn.hidden_neurons)
            hidden_neuron = 0.0;
            for j=1:length(inputs)
                hidden_neuron =hidden_neuron + inputs(j) *
nn.hidden_weights(j,i);
            end

```




Task 11.

Change the training data in `xor.m` to implement a different logical function, such as NOR or AND. Plot the error function of a successful trial.


NOR

```
training_set_input = [  
    0,0;  
    0,1;  
    1,0;  
    1,1  
];
```



AND

```
training_set_output = [  
    0;  
    0;  
    0;  
    1  
];
```



-----Thank You-----

ⁱ www.mind.ilstu.edu/curriculum/artificial_neural_net/xor_problem_and_solution.